

Table of Contents

1. Working up to Correlated Subqueries 1
2. Correlated Subqueries..... 3

1. Working up to Correlated Subqueries

We want to find out how many items (total quantity) are on each order we have. The first of these queries uses a scalar subquery in the Select but the output is not very interesting

Demo 01:

```
select customer_id, order_id
, (
    select sum(quantity_ordered)
    from oe_orderDetails OD
    ) as "NumItems"
from oe_orderHeaders OH
where rownum <=5;
```

CUSTOMER_ID	ORDER_ID	NumItems
409150	218	950
409150	223	950
409160	224	950
409160	225	950
403050	227	950

This seems to say that every order has a total quantity of 950 items; that does not look correct. The scalar subquery calculated the total number of items purchased on **all** orders. Look at the subquery. It says:

```
select sum(quantity_ordered)
from oe_orderDetails OD
```

What we want for each order is the total number of items **for this order**. We have to tie that subquery calculation to the Order id in the Order headers table. This is done via a correlation- the table in the subquery (order details) is tied to/joined to/correlated with the order header row we are looking at. This is different.

The subquery has a Where clause that correlates the subquery to the main query. OH is an alias for a table in the main query.

```
( select sum(quantity_ordered)
    from oe_orderDetails OD
    where OH.order_id = OD.order_id)
```

Demo 02: Now our result set makes more sense and is more useful

```
select customer_id, order_id
, ( select sum(quantity_ordered)
    from oe_orderDetails OD
    where OH.order_id = OD.order_id) as "NumItemsPerOrder"
from oe_orderHeaders OH
where customer_id IN ( 404100, 903000)
order by customer_id, order_id;
```

CUSTOMER_ID	ORDER_ID	NumItemsPerOrder
404100	303	1
404100	605	43
404100	2503	1
404100	2504	1
404100	2805	29
903000	306	2

903000	312	50
903000	313	1
903000	550	11
903000	551	20
903000	610	2
903000	2120	10
903000	2121	21
903000	3810	2

14 rows selected.

The nested sub query has a Where clause that refers to the outer query. That way we get the total of the quantity for this particular order. **This is a correlated subquery.** The table reference in the subquery is joined to the table references in the outer query.

The main query uses the order headers table and we get one row per order header row. We are not doing a group by clause.

Demo 03: Grouping with an outer join. It is not unusual to have more than one way to do a query.

```
select customer_id, OH.order_id ,sum(quantity_ordered) AS NumPerOrder
from oe_orderHeaders OH
left join oe_order_details OD on OH.order_id = OD.order_id
where customer_id IN ( 404100, 903000)
group by Customer_id, OH.order_id
order by customer_id, OH.order_id
;
```

Demo 04: We might want to get the number of items per order and the cost of the order. This uses two subqueries.

```
select customer_id, order_id
, ( select SUM (quantity_ordered)
  from oe_orderDetails OD
  where OH.order_id = OD.order_id) as "NumPerOrder"
, ( select SUM (quantity_ordered * quoted_price)
  from oe_orderDetails OD
  where OH.order_id = OD.order_id) as "OrderCost"
from oe_orderHeaders OH
order by customer_id,order_id
;
```

CUSTOMER_ID	ORDER_ID	NumPerOrder	OrderCost
400300	378	10	4500
401250	106	1	255.95
401250	113	1	22.5
401250	119	10	225
401250	301	1	205
401250	552	10	157.3
401250	2506		
401890	112	2	99.98
401890	519	6	114.74
402100	114	5	625
.			

This query returns rows where we have an Order header row with no Detail rows. Why does it do that? How would you change the query to return only Order header row that have Detail rows?

2. Correlated Subqueries

With a non-correlated subquery, the inner query could work on its own. With a correlated subquery, the inner query refers to attributes found in the outer query. That means that the correlated subquery cannot be run independently. Logically, the outer query works on the first row and processes the subquery using the attributes in the first row; then the outer query works on the second row and then processes the subquery using the attributes in the second row; it then continues through the rest of the rows in the outer query reevaluating the subquery repeatedly.

Some of the following are correlated subqueries. Although a correlated subquery may seem inefficient, the efficiency depends on the optimizer for the database engine.

If you have the choice of solving a task with a correlated query or with non-correlated query, you should generally choose the non-correlated version.

Demo 05: This uses an aggregate function to get the average price for all products.

```
select round( Avg ( prod_list_price),2)
from prd_products;
```

```
AVG(PROD_LIST_PRICE)
-----
          117.67
```

Demo 06: This uses a subquery to get products that cost more than the average price for all products.

```
select prod_id, prod_list_price, catg_id
from prd_products
where prod_list_price > (
    select round( Avg ( prod_list_price),2)
    from prd_products
);
```

PROD_ID	PROD_LIST_PRICE	CATG_I
1000	125.00	HW
1010	150.00	SPG
1090	149.99	HW
1040	349.95	SPG
1050	269.95	SPG
1060	255.95	SPG
1160	149.99	HW
4567	549.99	PET
4568	549.99	PET
4569	349.95	APL
1120	549.99	APL
1125	500.00	APL
1126	850.00	APL
1130	149.99	APL

Demo 07: This uses grouping and an aggregate function to get the average price for each product category.

```
select round( Avg ( prod_list_price),2) as AvgPrice, catg_id
from prd_products
group by catg_id;
```

```
AVGPRICE  CATG_ID
-----
123.17  PET
 23.88  HD
 13.88  MUS
```

```
67.64 HW
178.13 SPG
479.99 APL
8.75 GFD
```

Demo 08: We can use a **correlated subquery** to get the products that cost more than the average price for the **same type of products**. Notice that `prd_products` occurs in both the parent and the child query so we need to use table aliases in the join.

```
select prod_id, catg_id, prod_list_price
from prd_products Otr
where prod_list_price > (
    select Round( Avg ( prod_list_price),2)
    from prd_products Inr
    where Otr.catg_id = Inr.catg_id )
order by catg_id, prod_list_price
;
```

PROD_ID	CATG_I	PROD_LIST_PRICE
1125	APL	500.00
1120	APL	549.99
1126	APL	850.00
5000	GFD	12.50
5005	HD	45.00
1000	HW	125.00
1090	HW	149.99
1160	HW	149.99
2746	MUS	14.50
2747	MUS	14.50
2987	MUS	15.87
2984	MUS	15.87
2337	MUS	15.87
2234	MUS	15.88
2014	MUS	15.95
4567	PET	549.99
4568	PET	549.99
1060	SPG	255.95
1050	SPG	269.95
1040	SPG	349.95

20 rows selected.

Consider the subquery:

```
select AVG(prod_list_price)
from prd_products Inr
where Otr.catg_id = Inr.catg_id
```

This does not include a Group By clause. The subquery looks at only one value for `catg_id` - the one that matches the value for `catg_id` in the outer query for the current row being considered.

Since it is looking at only one category id, it will find only one average and so we can use the average in a filter of the type `Price > average`.

The one thing that should look odd about the subquery is that it refers to a table with an alias `Otr` that is not part of the subquery. That is where the "correlated" part of the correlated subquery comes in.

So let's go back to the main query. It gets one row from the product table and then tries to compare the price of that row to the average- what average; the average calculated by the subquery- which is the average for the same category id as that on the products table row we are looking at.

So the way to think of this query is

- working one row at a time through the products table
- for each row in the products table (one at a time) calculate the average price for that product id
- if the price for that product row is > average price for that category id, then return it.

This makes it sound like a very inefficient way to do this. Imagine we have a product tables of 50,000 rows of 10 different category ids. If the dbms actually carried the query out as I just described, that would mean calculated the average price 50,000 times. The dbms generally has a more efficient way - internally- to do this type of query. But logically you should think of the query as working with one row from the outer query processed against the subquery.

We want to find orders that are unusually high for a customer. Because we don't have a lot of rows, I defined this as an order that is more than 1.25 times the average order cost for that customer. This uses a CTE to assemble the data being used.

Demo 09: We will need the average order size by customer.

```
With OE_OrdExtTotal as (
  select OH.customer_id
    , order_id
    , sum ( OD.Quantity_ordered * OD.quoted_price) AS ordertotal
  from oe_orderHeaders OH
  join oe_orderDetails OD  using(order_id)
  group by OH.customer_id, order_id
)
select customer_id as "custid"
, to_char(avg(ordertotal), 9999.99) as "AvgPrice"
, to_char(1.25 * AVG(ordertotal), 9999.99) as "Cut_off"
from OE_OrdExtTotal
group by customer_id
order by customer_id;
```

custid	AvgPrice	Cut_off
400300	4500.00	5625.00
401250	173.15	216.44
401890	107.36	134.20
402100	1092.32	1365.40
403000	727.30	909.12
403010	1900.00	2375.00
403050	269.23	336.54
403100	218.84	273.55
. . . rows omitted		

Demo 10: Now we need a **correlated subquery** to compare a particular order with average orders for that customer.

```
With OE_OrdExtTotal as (
  select OH.customer_id
    , order_id
    , SUM ( OD.Quantity_ordered * OD.quoted_price) AS ordertotal
  from Oe_orderHeaders OH
  join Oe_orderDetails OD  using(order_id)
  group by OH.customer_id, order_id
)
```

```

select customer_id as "custid"
, order_id as "ordid"
, to_char( ordertotal, 9999.99) as "OrderCost"
from OE_OrdExtTotal OTR
where OrderTotal > 1.25 * (
    select AVG (ordertotal)
    from OE_OrdExtTotal INR
    where OTR.Customer_id = INR.Customer_id
    group by customer_id
)
order by customer_id;

```

custid	ordid	OrdTotal
401250	106	255.95
401250	119	225.00
402100	115	2305.00
403000	390	1400.00
403000	528	2629.00
403000	105	1205.40
403000	395	2925.00
403050	527	440.47
. . . rows omitted		

Let's start with the first customer shown here- customer_id 401250. This customer has 5 orders:

CUSTOMER_ID	ORDER_ID	ORDERTOTAL
401250	113	22.50
401250	552	157.30
401250	301	205.00
401250	119	225.00
401250	106	255.95

The cutoff for this customer is 216.44. Our query returns only two of this customers's orders- the ones where the order total is over this customer's cutoff.

401250	106	255.95
401250	119	225.00

Demo 11: To get customers with more than one order, we can use the count function for each customer_id as it occurs in the outer query; we do not need to qualify customer_id with Oe_order_headers in the inner query.

```

select customer_id
, customer_name_last
from cust_customers
where 1 < (
    select count ( *)
    from oe_orderHeaders
    where customer_id = cust_customers.customer_id )
;

```

CUSTOMER_ID	CUSTOMER_NAME_LAST
401250	Morse
401890	Northrep
402100	Morise
403000	Williams
403050	Hamilton
403100	Stevenson
. . . rows omitted	

Demo 12: Here the correlated subquery returns a number which is used as a parameter to a case expression. We want to use the number of orders for this customer- not for all customers.

```
select customer_id, customer_name_last
, Case (
    select count(*)
    from oe_orderHeaders OH
    where OH.customer_id = CS.customer_id
)
when 0 then '. . . No orders'
when 1 then '1 order'
when 2 then '2 orders'
when 3 then '3 orders'
else '4+ orders'
end as NumberOfOrders
from cust_customers CS
;
```

CUSTOMER_ID	CUSTOMER_NAME_LAST	NUMBEROFORDERS
401250	Morse	4+ orders
401890	Northrep	2 orders
402100	Morise	3 orders
402110	Coltrane	. . . No orders
402120	McCoy	. . . No orders
402500	Jones	. . . No orders
403000	Williams	4+ orders
403010	Otis	1 order
403050	Hamilton	4+ orders
403100	Stevenson	4+ orders
. . . rows omitted		