

## Table of Contents

1. Views .....	1
1.1. Reason for views .....	1
1.2. Example of creating and using a view .....	1
1.3. View versus CTE or subqueries .....	3
2. Some view syntax rules .....	3
3. Updating with Views .....	4
3.1. Determining which columns in a view can be changed .....	5
3.2. Updatable view using a join .....	6
4. Views that are more limiting.....	8
4.1. Declaring the view to be read-only .....	8
4.2. Making some columns not updatable.....	8
4.3. Updates with check constraints.....	9

## 1. Views

### 1.1. Reason for views

Suppose you have a fairly complex query dealing with customer orders that you need to run often. The query joins several tables and uses calculated columns etc. When you run this query, you filter for different order date ranges but the rest of the query is the same. You have asked the dba for help getting the query to be efficient and you do not want to keep writing the query over each time you need to run one of these reports. What you can do is create a database object called a view which encapsulates the complex query but still lets you add a filter for the order dates when you run that query. When you create your query, you can use the view in the From clause and not worry about the details of the join and calculated columns (not worrying about the details is a good thing!). You could also make the view available to programmers who may not have the experience to write complex queries.

A view is essentially a named, saved Select query. The view does not store any data; instead it stores the directions (the query) for getting the data. Once the view is created, you can use it as the table source for other queries. The definition of the view is stored on the server and is retrieved when the view is used. Each time you use the view, it goes to the actual base tables and retrieves the current data. The view metadata is stored in the data dictionary.

There are a few common reasons for using views

- to create a virtual table that presents a simplified set of column for queries. The view code might include joins and filters and calculated columns that the user of the view does not have to deal with directly
- to separate the logical database structure from the physical structure. The physical structure includes the objects create with the Create Table statement. The logical structure includes those tables and also the views. That way different users can have a different logical structure that reflects the user's needs.
- to improve security by letting a user see only certain columns and specified rows of a table.

### 1.2. Example of creating and using a view

Demo 01: We will use a fairly simple query for our view. When you work on a view, work out the details of the select statement first.

```
create or replace view custReportGoodCredit_01 as
select customer_id as CustomerID
, customer_name_first || ' ' || customer_name_last as CustomerName
from cust_customers
where customer_credit_limit > 3000
with read only;
```

**Demo 02: Using the view**

```
select CustomerName, CustomerID
from custReportGoodCredit_01
order by CustomerID;
```

CUSTOMERNAME	CUSTOMERID
-----	-----
Arnold McGold	400300
Sally Williams	403000
Elisha Otis	403010
Alexis Hamilton	403050
James Stevenson	403100
... rows omitted	

Note that the person using this view cannot see the credit limit value because it is not exposed by the view.

Data displayed through a view will reflect the current data in the table at the time the query was run; not the data at the time the view was defined. A person using a view as the table expression might not be aware that this is a view and not a base table.

**Demo 03: Creating a view that concatenates the customer names and calculates the linetotal**

```
create or replace view ordreport_01 as
select
  order_id      as orderid
,  order_date   as orderdate
,  customer_id  as customerid
,  customer_name_first || ' ' || customer_name_last as customername
,  prod_id      as itempurchased
,  prod_name    as itemdescription
,  quoted_price * quantity_ordered as linetotal
from cust_customers
join oe_orderHeaders using (customer_id)
join oe_orderDetails using (order_id)
join prd_products using (prod_id)
where quoted_price > 0 and quantity_ordered > 0 with read only;
```

**Demo 04: Describing the view.** One thing to check is that the data types of the columns in the view are appropriate. Do not over format the columns you wish to filter; do not output a formatted number column if you might wish to filter on it. Remember that when you format a numeric column with To\_Char you are returning a string- not a number.

```
Desc ordReport_01
```

Name	Null?	Type
-----	-----	-----
ORDERID	NOT NULL	NUMBER(38)
ORDERDATE	NOT NULL	DATE
CUSTOMERID	NOT NULL	NUMBER(38)
CUSTOMERNAME		VARCHAR2(51)
ITEMPURCHASED	NOT NULL	NUMBER(38)
ITEMDESCRIPTION	NOT NULL	VARCHAR2(25)
LINETOTAL		NUMBER

**Demo 05: Using the view this is a lot easier than repeating that join for each query**

```
select orderid, orderdate, itempurchased, linetotal
from ordreport 01;
```

ORDERID	ORDERDATE	ITEMPURCHASED	LINETOTAL
-----	-----	-----	-----
105	01-OCT-15	1020	155.4

105	01-OCT-15	1030	300
105	01-OCT-15	1010	750
106	01-OCT-15	1060	255.95
107	02-OCT-15	1110	49.99
. . .	rows omitted		

**Demo 06: Using the view with a filter**

```

Select
 orderid
, orderdate
, itempurchased
From ordreport_01
Where itempurchased in ( 1071, 1072)
order by orderdate;

```

ORDERID	ORDERDATE	ITEMPURCHASED
411	01-AUG-15	1071
411	01-AUG-15	1072
408	20-NOV-15	1071
2122	23-JAN-16	1071
2122	23-JAN-16	1072
4510	01-FEB-16	1071
4510	01-FEB-16	1072

**Demo 07: Alternate syntax when creating a view- this defines the column names in the view header rather than in the select**

```

create or replace view ordreport_02(orderid, orderdate, customerid, linetotal)
as
select order_id, order_date, customer_id
, quoted_price * quantity_ordered
from oe_orderheaders
join oe_orderdetails using(order_id)
join prd_products using(prod_id)
where quoted_price > 0 and quantity_ordered > 0
order by order_id
with read only;

```

**1.3. View versus CTE or subqueries**

You may have noticed a similarity between CTE and views. A view is stored in the database as a persistent database object; a CTE exists for the lifetime of the query. So if you have a result set that you want recalculated each time you use it and you would use that result set often or you want other users to use that result set, then a View is a good idea. If you want to break a complex task into pieces to solve one step at a time and the task is unique and it is not sometime you want for other tasks, then a CTE is a good idea.

A production database will often have a large number of views set up for business purposes but you don't want to create a view that is used only once.

**2. Some view syntax rules**

A view is an object created in a specific database.

- You can use either the Create View syntax or the Create or Replace View syntax. If the name you select for the view is already used by another view, the Create view syntax will report an error; the Create or Replace View syntax will replace the current definition of the view.

- You cannot have a view with the same name as a table.
- You cannot create a view on a table that does not exist yet unless you include the word Force. You cannot use the view until the table is created, but this keyword will let you define the view.
- Since a view is a stored object, you can drop the view with a drop view command.  
`drop view my_view;`
- Your user account has to have permissions to create views. Your CCSF Oracle account has that permission.
- You can assign new names for the columns when you define a view or you can use the names in the underlying tables.
- If you have a calculated column in the underlying select, you must provide an alias for that column.
- If you define the columns in the view header, the number of columns in the view definition must match the number of columns in the select used to define the view.
- When you create an alias within a view, the alias should be a single word and not enclosed in quotes
- Do not create a view using Select \*; the view will be stored with the \* expanded into the column names at the time the view was created and will not be updated if the underlying tables are changed.
- Oracle allows an Order By clause in the definition of a view

### 3. Updating with Views

Some restrictions for updating (any change) through a view. We will discuss some of these options in this document. If you think about these rules and the need for the dbms to understand exactly what data to change, these rules make more sense.

- If the view is based on all of the columns of a single table and the view was not created with a read-only option or a with check option, then you can update through the view.
  - A read only option means that the view cannot be used for updates
  - A With Check option means that you cannot do an update that would violate the check option.
  - You cannot update a column in a view which is based on a calculated column. If the view has several columns and only some columns are calculated, you can update the other columns
  - If you have any calculated columns, then you cannot insert via the view. However you might be able to delete via the view.
  - You cannot update views that are defined with aggregate functions or Group By
  - You cannot update views that are defined with Distinct or rownum
  - You cannot update a view created with a set operation such as Union
- 
- If the view is based on multiple tables, you cannot do an update on a column which maps to a non key-preserved table. (Essentially this means that if the view definition does not contain the pk – as a pk for the view set, then the view cannot be updated. So if we had a parent and child situation and the view definition includes the pk of the parent we would, in general, be able to update the parent. But we could not update the child even with the child table pk since it could occur more than once in the view rowset)

Demo 08: In order to not change the data in the regular table, insert the following rows for testing.

An Insert for the employees table

```
delete from emp_employees where emp_id = 1995;
insert into emp_employees values
(1995, 'Zahn', 'Joe', '111111111', 100, 10, date'2009-09-09', 500, 2);
```

An Insert for the products table

```
delete from prd_products where prod_id = 1995;
insert into prd_products values
(1995, 'book', 'Train your cat book', 29.95, 'PET');
```

### 3.1. Determining which columns in a view can be changed

Demo 09: What is the structure of emp\_employees?

```
desc emp_employees
```

Name	Null?	Type
EMP_ID	NOT NULL	NUMBER(38)
NAME_LAST	NOT NULL	VARCHAR2(25)
NAME_FIRST		VARCHAR2(25)
SSN	NOT NULL	CHAR(9)
EMP_MNG		NUMBER(38)
DEPT_ID	NOT NULL	NUMBER(38)
HIRE_DATE	NOT NULL	DATE
SALARY		NUMBER(8,2)
JOB_ID	NOT NULL	NUMBER(38)

Demo 10: Simple, single table view. This view can be used for updates of the employee name and the id

```
CREATE OR REPLACE VIEW vw_Emp AS
select emp_id, name_last, name_first
from emp_employees;
```

Demo 11: Which columns can be changed? Use the data dictionary view user\_updatable\_columns

```
select column_name, updatable
from user_updatable_columns
where table_name = 'VW_EMP';
```

COLUMN_NAME	UPDATABLE
EMP_ID	YES
NAME_LAST	YES
NAME_FIRST	YES

**Discussion:** This update works.

```
Update vw_emp set name_first = 'Susan' where emp_id = 1995;
```

This one does not work; the attribute salary is not accessible through the view.

```
Update vw_emp set salary = 12500 where emp_id = 1995;
```

```
ORA-00904: "SALARY": invalid identifier
```

This update does not work; the attribute name\_last is accessible through the view, but that attribute is not nullable in the base table. The table constraints must always be met.

```
Update vw_emp set name_last = null where emp_id = 1995;
```

```
ORA-01407: cannot update ("ROSE151A"."EMP_EMPLOYEES"."NAME_LAST") to NULL
```

Can you do an insert via this view? No- there are several attributes in the table which are not nullable and not accessible via the view.

Can you do a delete via this view? Yes, unless there is another constraint that would disallow the delete.

```
Delete from vw_emp where emp_id = 1995;
```

If you have removed or changed that row, reinsert it.

Demo 12: This is a view that shows summary data

```
CREATE OR REPLACE VIEW vw_HighPriceByCategory AS
select Catg_id, MAX(prod_list_price) As HighPrice
from prd_products
group by Catg_id;
```

Which columns can be changed? You cannot change data in a view with a group by.

```
select column_name, updatable
from user_updatable_columns
where table_name = 'VW HIGHPRICEBYCATEGORY';
```

COLUMN_NAME	UPDATABLE
CATG_ID	NO
HIGHPRICE	NO

If we try to change data through the view, we get an error because the attribute HighPrice does not tie back to a specific row in the underlying table

```
update vw_HighPriceByCategory
set HighPrice = 250
where Catg_id = 'PET';
```

ORA-01732: data manipulation operation not legal on this view

### Demo 13: This is a view that includes a calculated column

```
CREATE OR REPLACE VIEW vw_NewPrice AS
select Catg_id, prod_name, Round(prod_list_price * 1.05,2) As NewPrice
from prd_products;
```

Some of these columns can be updated but we cannot update the calculated column.

```
select column_name, updatable
from user_updatable_columns
where table_name = 'VW NEWPRICE';
```

COLUMN_NAME	UPDATABLE
CATG_ID	YES
PROD_NAME	YES
NEWPRICE	NO

Trying to update the calculated column is blocked.

```
Update vw_NewPrice
set NewPrice = 45
where prod name = 'Train your cat book';
```

ORA-01733: virtual column not allowed here

### Demo 14: This is another view that includes a calculated column

```
CREATE OR REPLACE VIEW vw_NewPrice2 AS
select Catg_id, prod_name, prod_list_price * 1 As NewPrice
from prd_products;
```

If you think about this mathematically,  $\text{prod\_list\_price} * 1$  does not really change the value, but this is still a calculated column and cannot be updated via the view.

```
Update vw_NewPrice2
set NewPrice = 45
where prod name = 'Train your cat book';
```

ORA-01733: virtual column not allowed here

## 3.2. Updatable view using a join

Create a view that joins the department table (the parent) and the employee table (the child). Can we update columns in the parent? in the child? There are several different ways to do the join. The next three views use different join techniques

**Demo 15: Creating the join with the USING clause.**

```
CREATE OR REPLACE VIEW vw_em_dept_1 AS
  select  emp_id, name_last
        ,    dept_id
        ,    dept_name
  from    emp_employees join emp_departments using(dept_id);
```

**Demo 16: Creating the join with the ON clause including the dept\_id column from the child table.**

```
CREATE OR REPLACE VIEW vw_em_dept_2 AS
  select  emp_id, name_last
        ,    emp_employees.dept_id
        ,    dept_name
  from    emp_employees
  join    emp_departments
  on      emp_employees.dept_id = emp_departments.dept_id;
```

**Demo 17: Creating the join with the ON clause including the dept\_id column from the parent table.**

```
CREATE OR REPLACE VIEW vw_em_dept_3 AS
  select  emp_id, name_last
        ,    emp_departments.dept_id
        ,    dept_name
  from    emp_employees
  join    emp_departments
  on      emp_employees.dept_id = emp_departments.dept_id;
```

If you do a Select \* for any of these three views you get the same result set.

**Demo 18: But there is a difference. Which columns are updatable ? (The Break command is used to make the output easier to read)**

```
BREAK ON table_name SKIP 1

select table_name, column_name, updatable
from user_updatable_columns
where table_name LIKE 'VW EM DEPT %' ;
```

TABLE_NAME	COLUMN_NAME	UPDATABLE
VW_EM_DEPT_1	EMP_ID	YES
	NAME_LAST	YES
	DEPT_ID	NO
	DEPT_NAME	NO
VW_EM_DEPT_2	EMP_ID	YES
	NAME_LAST	YES
	DEPT_ID	YES
	DEPT_NAME	NO
VW_EM_DEPT_3	EMP_ID	YES
	NAME_LAST	YES
	DEPT_ID	NO
	DEPT_NAME	NO

(This is rather complex- I mostly want you to understand that if a view is based on a join of tables, updating gets more complex.)

In all of these we can update the Emp\_id and name\_last which come from the child table. For each Emp\_id or Emp\_name, that value ties back to a single definable row in the employees table. There is only one row in the employees table for each row in the view results table. ( the term "key-preserved" is used for a base table that

has a 1-1 relationship with the rows in the views- therefore a row in the view can be unambiguously tied to a row in that base table.) We cannot update the dept\_name which comes from the parent table; there are multiple rows in the result set for a row in the departments table.

The dept\_id can be updated in the second view- where that attribute in the view comes uniquely from the child table. The dept\_id cannot be updated in the third view- where that attribute in the view comes from the parent table and is not unique. The dept\_id cannot be updated in the first view- where that attribute in the view comes from the new generated column and not from either table.

Demo 19: An attempt to update the dept\_id using the third view. This fails.

```
update vw_em_dept_3
set dept_id = 80
where emp_id = 1995;
```

```
set      dept_id = 80
      *
```

ERROR at line 2:  
ORA-01779: cannot modify a column which maps to a non key-preserved table

**Discussion:** We can do the update via the second view. In that case we are changing the dept id for child table- (the employee table). So this changes a single row. If we try to do the change via the third view, we would be changing the dept\_id in the department table- so presumably we would be changing dept\_id 10 to 80 in the department table- that would result in two rows in the department table with ID 80. Allowing this type of change would create problems. ( Again, if you did the update, do a rollback.)

## 4. Views that are more limiting

There are some options that can be used in the creation of a view that provide protection of the data in the base table.

### 4.1. Declaring the view to be read-only

Demo 20: Declaring that the view is read only. None of the attributes are updatable.

```
CREATE OR REPLACE VIEW ProductList AS
select prod_id, prod_list_price, prod_name, prod_desc
from prd_products
WITH READ ONLY;
```

### 4.2. Making some columns not updatable

Demo 21: The list price and the description can be updated but not the id or name. because these are calculated columns

```
CREATE OR REPLACE VIEW ProductList2 AS
select  prod_id + 0 as prod_id
        , prod_list_price as prod_list_price
        , prod_name || ' ' as prod_name
        , prod_desc
from    prd_products;

select column_name, updatable
from user_updatable_columns
where table_name = 'PRODUCTLIST2';
```

COLUMN_NAME	UPD
PROD_ID	NO
PROD_LIST_PRICE	YES



PROD_NAME	NO
PROD_DESC	YES

### 4.3. Updates with check constraints

The next few SQL statements show the purpose of the With Check Option clause. Suppose I create a view that has a WHERE clause. I can use the view to update the data in such a way that the data no longer satisfies that WHERE clause. The With Check Option clause restricts this behavior.

#### Demo 22: Creating a view with a check constraint

```
CREATE OR REPLACE VIEW EmpDept80 AS
  select   emp_id, name_last
    ,      dept_id
  from     emp_employees
  where    dept_id = 80
WITH CHECK OPTION;
```

#### Which columns can be changed?

COLUMN_NAME	UPD
-----	---
EMP_ID	YES
NAME_LAST	YES
DEPT_ID	YES

#### Demo 23: It is OK to change the last name value through this view. (Be certain that you have employee 1995 in dept 80 before you run this one.)

```
update   EmpDept80
set       name_last = 'Hiller'
where     emp_id = 1995;
```

#### Demo 24: It is not OK to change the department to a value that does not meet the check constraint.

```
update   EmpDept80
set       dept_id = 20
where     emp_id = 1995;
```

```
ORA-01402: view WITH CHECK OPTION WHERE -clause violation
```

#### Demo 25: Creating a view with a check constraint

```
CREATE OR REPLACE VIEW EmpHighSalary AS
  select   emp_id, name_last, salary
    ,      dept_id
  from     emp_employees
  where    salary > 70000
WITH CHECK OPTION;
```

```
-- the record set
select *
from EmpHighSalary;
```

EMP_ID	NAME_LAST	SALARY	DEPT_ID
-----	-----	-----	-----
100	King	100000	10
101	Koch	98005	30
205	Higgs	75000	30
162	Holme	98000	35
146	Partne	88954	215
161	Dewal	120000	215

206 Geitz	88954	30
204 King	99090	30

8 rows selected.

Demo 26: We can update employee 1995 salary via the employees table and then that employee will appear via the view.

```
update emp_employees
set name_last = 'Prince',
    salary = 75000
where emp_id = 1995;
```

```
select *
from EmpHighSalary ;
```

EMP_ID	NAME_LAST	SALARY	DEPT_ID
100	King	100000	10
101	Koch	98005	30
205	Higgs	75000	30
162	Holme	98000	35
146	Partne	88954	215
161	Dewal	120000	215
206	Geitz	88954	30
204	King	99090	30
1995	Prince	75000	10

9 rows selected.

Demo 27: Updating through the view. This will do the update.

```
update EmpHighSalary
set name_last = 'Prince',
    salary = 85000
where emp_id = 1995;
```

1 row updated.

Demo 28: Updating through the view. This update will fail due to the check constraint

```
update EmpHighSalary
set name_last = 'Koch',
    salary = 10000
where emp_id = 1995;
```

ORA-01402: view WITH CHECK OPTION where-clause violation

Similarly you cannot do an Insert via a view with a check constraint if that insert would not be displayed with the view- i.e. if the insert data does not meet the Where clause criterion.

It would be a good idea to reload the Altgeld tables at this time.