

**Table of Contents**

1. Aggregate () over ().....	1
1.1. Partition.....	2
1.2. ListAgg .....	4
2. Running Totals .....	5
3. Windowing clause.....	6
3.1. The Range phrase .....	7
3.2. Moving Windowing Clause .....	8
3.3. Logical Windowing.....	11

**1. Aggregate () over ()**

Suppose you want to display each person's salary and how much their salary is over the average salary for all employees.

We'll start by looking just at department 20. We have four employees. The sum of their salaries is 81000 ; the average salary (rounded to an integer) is 20250 . We could do this with a subquery in the Select.

```
variable dpt number
exec :dpt := 20;
```

**Demo 01: Using a subquery and the avg function**

```
select emp_id, salary
, Round(salary -
      (select Avg(salary)
       from adv_emp
       where dept_id = :dpt),2) as Over_under_avg
from adv_emp
where dept_id = :dpt
;
```

EMP_ID	SALARY	OVER_UNDER_AVG
302	14000	-6250
303	27000	6750
312	28000	7750
315	12000	-8250

**Demo 02: You could also use a CTE and a Cross join. Be sure you understand why a cross join will work here**

```
with avgSal as (
  select Avg(salary * 1.0) as AvgDept100
  from adv_emp
  where dept_id = :dpt
)
select emp_id, salary
, Round(salary - AvgDept100,2) as Over_under_avg
from adv_emp
cross join avgSal
where dept_id = :dpt;
```

**Demo 03: Using avg() Over()**

We can also do this with an Avg() Over function to get the same result. This is a much simpler syntax than the subquery. If we just do the *avg(salary) over()* for the second column, you can see that the average is calculated over dept\_id 20 for each row.

So we do the subtraction to get the offset of each employee's salary compared to the average.

```
select emp_id, salary
, Round(salary - ( Avg(salary) Over() ), 2 ) as Over_under_avg
from adv_emp
where dept_id = :dpt;
```

In this case the Over() clause has no argument; this is referred to as a Null Over clause and it means that the function applies to the entire dataset- since the function is calculated after the Where clause, this means to the rows for dept\_id 20 only. In the subquery version we had to do the filter in both the subquery and the parent query to get the correct result.

### 1.1. Partition

We might want to look at all the employees and check their over\_under\_avg based on their dept\_id only. This means we want to group the employees by dept\_id and calculate the average for each group separately. That is a partition. This query will give that result. I have sorted by Dept ID and salary to make the output easier to read.

#### Demo 04: Using avg() Over() with a partition

```
select dept_id, salary
, Round(salary - ( Avg(salary) Over( Partition by dept_id ) ), 0 )
as Over_under_avg
, emp_id
from adv_emp
order by dept_id, emp_id
;
```

DEPT_ID	SALARY	OVER_UNDER_AVG	EMP_I
10	15000	-11111	301
10	27000	889	305
10	30000	3889	309
10	25000	-1111	310
10	28000	1889	311
10	30000	3889	319
10	25000	-1111	320
10	30000	3889	321
10	25000	-1111	322
15	25000	0	323
20	14000	-6250	302
20	27000	6750	303
20	28000	7750	312
20	12000	-8250	315
30	28000	7111	304
30	28000	7111	306
30	13500	-7389	307
30	15000	-5889	308
30	11000	-9889	313
30	30000	9111	314
30	26000	5111	316
30	25000	4111	317
30	11500	-9389	318

23 rows selected.

**Demo 05: Who earned more than the average salary for their department?**

```

with Aggs as (
  select emp_id, dept_id, salary
    , Avg(salary) Over( Partition by dept_id) as dept_avg
  from   adv_emp
)
select emp_id, dept_id, salary
from aggs
where salary > dept_avg
order by dept_id, emp_id;

```

EMP_I	DEPT_ID	SALARY
305	10	27000
309	10	30000
311	10	28000
319	10	30000
321	10	30000
303	20	27000
312	20	28000
304	30	28000
306	30	28000
314	30	30000
316	30	26000
317	30	25000

12 rows selected.

**Demo 06: Calculating Percent of total: Now we want to know what each employee's salary is as a percent of the total salary for that department.**

```

select dept_id, emp_id, salary
  , round(salary /(sum(salary) over (partition by dept_id) ) * 100, 2)
    as percent_dept_salary
from   adv_emp
order by dept_id, salary;

```

DEPT_ID	EMP_I	SALARY	PERCENT_DEPT_SALARY
10	301	15000	6.38
10	320	25000	10.64
10	310	25000	10.64
10	322	25000	10.64
10	305	27000	11.49
10	311	28000	11.91
10	319	30000	12.77
10	321	30000	12.77
10	309	30000	12.77
15	323	25000	100
20	315	12000	14.81
20	302	14000	17.28
20	303	27000	33.33
20	312	28000	34.57
30	313	11000	5.85
30	318	11500	6.12
30	307	13500	7.18
30	308	15000	7.98
30	317	25000	13.3
30	316	26000	13.83
30	306	28000	14.89
30	304	28000	14.89
30	314	30000	15.96

23 rows selected.

Start by looking at the results for dept 15. There is one employee, with a salary of 25000. This row reports as 100% of the department salary total. Then look at the results for dept 20. There are four employees. The total salary for dept 210 is 28000. Employee 315 has a salary of 12000 which is about 15% of the department total salary.

You could get the same result with the **ratio\_to\_report** function:

```
ratio_to_report ( salary ) over ( partition by dept_id )
```

#### Demo 07: Using Ratio\_to\_Report

```
select dept_id, emp_id, salary
, round(ratio_to_report ( salary ) over ( partition by dept_id ) * 100,2)
  as percent_dept_salary
from   adv_emp
order by dept_id, salary;
```

### 1.2. ListAgg

This is another aggregate function you can use with the analytical techniques. ListAgg returns a concatenated set of values.

Demo 08: Simple aggregate of the names with a semicolon followed by a space as the delimiter, We get a single row returned with the employees last names.

```
select listagg(name_last, ';' ' ')
      within group (order by name_last) as "Employee List"
from adv_emp;
```

Employee List

```
-----
Battaglia; Beiderbecke; Brubeck; Cohen; Coltrane; Davis; Ellington; Evans; Green; Hancock;
Jarrett; Mobley; Monk; Montgomery; Quebec; Redman; Rollins; Shorter; Tatum; Turrentine;
Wabich; Wabich; Wasliewski
```

Demo 09: We can order the values by salary in descending order

```
select listagg(name_last, ';' ' ')
      within group (order by salary desc) as "Employee List"
from adv_emp;
```

Employee List

```
-----
Beiderbecke; Redman; Rollins; Turrentine; Brubeck; Cohen; Ellington; Mobley; Coltrane;
Quebec; Monk; Jarrett; Montgomery; Wabich; Wabich; Wasliewski; Evans; Green; Hancock;
Tatum; Battaglia; Shorter; Davis
```

Demo 10: We can add a regular Group by clause and listagg returns aggregates for each group

```
select dept_id as "Dept",
      listagg(name_last, ';' ' ')
        within group (order by year_hired) as "Employees by Year Hired"
from   adv_emp
group by dept_id;
```

Dept Employees by Year Hired

```
-----
10 Green; Brubeck; Coltrane; Jarrett; Wabich; Wabich; Beiderbecke; Redman; Rollins
15 Montgomery
20 Ellington; Hancock; Battaglia; Quebec
```

```
30 Cohen; Mobley; Davis; Tatum; Evans; Monk; Turrentine; Shorter; Wasliewski 10 King
4 rows selected.
```

You can also put the dept\_id into a partition by clause in the function. Also try this without the Distinct keyword.

```
Select Distinct dept_id As "Dept"
, listagg(name_last, '; ')
  WITHIN GROUP
  (ORDER BY year_hired)
  OVER (PARTITION BY dept_id) as "Employees by Year Hired"
from adv_emp;
```

```
Dept Employees by Year Hired
-----
15 Montgomery
30 Cohen; Mobley; Davis; Tatum; Turrentine; Evans; Monk; Wasliewski; Shorter
20 Hancock; Ellington; Battaglia; Quebec
10 Green; Coltrane; Brubeck; Jarrett; Wabich; Wabich; Redman; Rollins; Beiderbecke
4 rows selected.
```

## 2. Running Totals

We can calculate the total of the salaries and get one value for the rows in the table.

Demo 11: Simple aggregate functions over the table

```
select sum(salary), count(*)
from adv_emp;
```

```
SUM(SALARY)    COUNT(*)
-----
529000         23
```

Suppose we want a running total instead; we want the first row to show the total of the first salary; the second row to show the total of row 1 to row 2; the third row to show the total of row 1 to row 3, etc.

The following query does that; it sums the salary over a range consisting of all of the rows from the start (**unbounded preceding**) to the **current** row. We want the rows in the running total to be in emp\_id order.

Demo 12: Running total over the table using Range Between

```
select emp_id, salary
, sum(salary) over(
    order by emp_id
    range between unbounded preceding and current row)
  as run_tot
from adv_emp
order by emp_id
;
```

```
EMP_ID  SALARY  RUN_TOT
-----
301      15000    15000
302      14000    29000
303      27000    56000
304      28000    84000
305      27000   111000
```

306	28000	139000
307	13500	152500
308	15000	167500
309	30000	197500
310	25000	222500
311	28000	250500
312	28000	278500
313	11000	289500
314	30000	319500
315	12000	331500
316	26000	357500
317	25000	382500
318	11500	394000
319	30000	424000
320	25000	449000
321	30000	479000
322	25000	504000
323	25000	529000

23 rows selected.

Be certain to have the two Order By clauses to sort on the same columns for the output to be understandable.

### 3. Windowing clause

By adding **Partition BY department\_id**, we have told SQL to do a running total for each department, restarting the total when the department changes. We are ordering by dept\_id, salary.

This is an example of a windowing frame- a set of rows creating a windows partition. The windowing frame is a moving frame. In this example the windowing frame consists of all rows from the last dept\_id change ( the partition by argument) up to and including the current row.

Demo 13: Running total by department, adding a partition clause.

```
select dept_id
, emp_id
, salary
, sum(salary) over(
    PARTITION BY DEPT_ID
    order by dept_id, salary
    RANGE BETWEEN UNBOUNDED PRECEDING AND CURRENT ROW)
    as run_tot
from adv_emp
order by dept_id, salary
;
```

I artificially added blank lines for each dept id break

DEPT_ID	EMP_ID	SALARY	RUN_TOT
-----			
10	301	15000	15000
10	320	25000	90000
10	310	25000	90000
10	322	25000	90000
10	305	27000	117000
10	311	28000	145000
10	319	30000	235000
10	321	30000	235000
10	309	30000	235000
15	323	25000	25000

20	315	12000	12000
20	302	14000	26000
20	303	27000	53000
20	312	28000	81000
30	313	11000	11000
30	318	11500	22500
30	307	13500	36000
30	308	15000	51000
30	317	25000	76000
30	316	26000	102000
30	306	28000	158000
30	304	28000	158000
30	314	30000	188000

### 3.1. The Range phrase

The Range phrase is called the Windowing-clause and it defaults to range between unbounded preceding and current row.

Demo 14: This query uses the default and gives us the same output as above.

```
select dept_id, emp_id, salary
, sum(salary) over(
    PARTITION BY DEPT_ID
    order by dept_id, salary
)
    as run_tot
from adv_emp
order by dept_id, salary;
```

Demo 15: Using the default range between unbounded preceding and current row and ordering by employee id

```
select emp_id, salary
, sum(salary) over(
    order by emp_id
    -- range between unbounded preceding and current row
)
    as run_tot
from adv_emp
order by emp_id;
```

EMP_I	SALARY	RUN_TOT
301	15000	15000
302	14000	29000
303	27000	56000
304	28000	84000
305	27000	111000
306	28000	139000
307	13500	152500
308	15000	167500
309	30000	197500
310	25000	222500
311	28000	250500
312	28000	278500
313	11000	289500
314	30000	319500
315	12000	331500
316	26000	357500

317	25000	382500
318	11500	394000
319	30000	424000
320	25000	449000
321	30000	479000
322	25000	504000
323	25000	529000

### 3.2. Moving Windowing Clause

The windowing-clause allows you to set a moving window of rows over which the function should be applied.

Suppose we want the current row and two rows before and after the current row. In the picture below the current row is highlighted in yellow and the window includes the blue and yellow highlighted rows

22	300	2006	70000
23	300	2006	80500
24	300	2006	80500
25	300	2006	6500
26	300	2007	80000
27	300	2007	8500
28	300	2007	50000
29	300	2007	65000
30	300	2008	60000
31	300	2008	65000
32	300	2008	50000
33	300	2008	80000

This shows the moving window frame. Note that at the start or the end of the dataset, the window will include less than 5 rows.

22	300	2006	70000
23	300	2006	80500
24	300	2006	80500
25	300	2006	6500
26	300	2007	80000
27	300	2007	8500
28	300	2007	50000
29	300	2007	65000
30	300	2008	60000
31	300	2008	65000
32	300	2008	50000

22	300	2006	70000
23	300	2006	80500
24	300	2006	80500
25	300	2006	6500
26	300	2007	80000
27	300	2007	8500
28	300	2007	50000
29	300	2007	65000
30	300	2008	60000
31	300	2008	65000
32	300	2008	50000

22	300	2006	70000
23	300	2006	80500
24	300	2006	80500
25	300	2006	6500
26	300	2007	80000
27	300	2007	8500
28	300	2007	50000
29	300	2007	65000
30	300	2008	60000
31	300	2008	65000
32	300	2008	50000

22	300	2006	70000
23	300	2006	80500
24	300	2006	80500
25	300	2006	6500
26	300	2007	80000
27	300	2007	8500
28	300	2007	50000
29	300	2007	65000
30	300	2008	60000
31	300	2008	65000
32	300	2008	50000


22	300	2006	70000
23	300	2006	80500
24	300	2006	80500
25	300	2006	6500
26	300	2007	80000
27	300	2007	8500
28	300	2007	50000
29	300	2007	65000
30	300	2008	60000
31	300	2008	65000
32	300	2008	50000

We can use several ways to express the range that compose a window. We can use a physical number of rows to be included or a term such as unbounded preceding. The windowing clause is often used when we have a series of values and we want to find a moving average. A moving average can be used with a time dependent series of data and we want to smooth out the data by looking at a three day average. That way if we have a few days that are somewhat out of the normal range they will be blended into the moving average. An example will help with this.

Suppose you were tracking sales over a period of several days. The sales values are apt to fluctuate each day. We could produce the graph of the sales shown on the left which shows each day's sales and we see some pretty big changes. We could also graph a three day average as shown on the right which smoothes out those one day changes. We still see variation in data but a single value does not show up as being that significant.

This is not a case of which graph is correct- but rather what do we want to see.





The test table is `adv_sales` with the first column being a number from 1-20 representing the day of the time period and the second column will be the sales for that day. The inserts are provided on the SQL included with these notes(demo 01). We can refer to this as a three-day average if we have a row, in the proper order, for each day from the first to the last - with no gaps or duplicates.

```
Create table adv_sales (
    sales_day number(2) primary key
    , sales number (5) check (sales >= 0)
);
```

The window we will use include the current row and the 2 preceding rows.

**Demo 16:** Here the window is the current row and the 2 preceding rows. The Column statement is an SQL\*Plus command to format the third column; it does not work in SQL Developer.

```
Column Three_day_avg format "9999.99"
```

```
select sales_day
, sales
, Avg(sales) OVER (
    ORDER BY sales_day
    ROWS BETWEEN
    2 PRECEDING AND CURRENT ROW
) AS three_day_avg
from adv_sales
order by sales day;
```

SALES_DAY	SALES	THREE_DAY_AVG
25-APR-15	400	400.00
26-APR-15	400	400.00
27-APR-15	400	400.00
28-APR-15	300	366.67
29-APR-15	900	533.33
30-APR-15	580	593.33
01-MAY-15	425	635.00
02-MAY-15	10	338.33
03-MAY-15	325	253.33
04-MAY-15	500	278.33
05-MAY-15	550	458.33
06-MAY-15	0	350.00
07-MAY-15	0	183.33
08-MAY-15	200	66.67
09-MAY-15	450	216.67
10-MAY-15	580	410.00
11-MAY-15	425	485.00
12-MAY-15	475	493.33
13-MAY-15	375	425.00
14-MAY-15	500	450.00
15-MAY-15	650	508.33

16-MAY-15	550	566.67
17-MAY-15	450	550.00
18-MAY-15	500	500.00
19-MAY-15	575	508.33
20-MAY-15	450	508.33
21-MAY-15	500	508.33
22-MAY-15	575	508.33
23-MAY-15	850	641.67
24-MAY-15	500	641.67
25-MAY-15	575	641.67
26-MAY-15	500	525.00
27-MAY-15	575	550.00
28-MAY-15	500	525.00
29-MAY-15	575	550.00
30-MAY-15	575	550.00
31-MAY-15	575	575.00
01-JUN-15	425	525.00
02-JUN-15	500	500.00
03-JUN-15	455	460.00
04-JUN-15	0	318.33
05-JUN-15	0	151.67
06-JUN-15	0	0.00
07-JUN-15	0	0.00
08-JUN-15	900	300.00
09-JUN-15	450	450.00
10-JUN-15	780	710.00
11-JUN-15	475	568.33
12-JUN-15	875	710.00
13-JUN-15	375	575.00
14-JUN-15	800	683.33

We can see that the last column has less variation since each value (except for the end points depending on row preceding or rows following) represents the average of three data points. The low number and the high number do affect the three\_Day\_avg, but their effect is not as much as in the sales column.

The first two rows are not three row averages, since we do not yet have three rows of sale. What Oracle does is treat the missing rows as Nulls. We would need to consider if this is relevant to the purpose for which we are running the query.

It may be that our company sells more merchandise at the start or the end of the month. In that case, the three-day-average could be misleading. If we were storing data about weather and the data stored the high temperature for each day, we might not think that that is influenced by the start or end of the month. Queries like this are generally done with large amounts of data and often the end points can be ignored or removed.

However you always need to check these assumptions. Maybe some factory in the area does larger runs on the last two days of the month and this could influence the local temperature due to atmospheric conditions. Assumptions in statistics can hide the data you are trying to find.

**Demo 17:** Here the window extends over four rows, the current, two days preceding and one day following. It also limits the display to some of the rows in the table.

---

Column Four\_day\_avg format "9999.99"

```
select sales_day
, sales
, round(avg(sales) over (
    order by sales_day
    rows between 2 preceding and 1 following ), 2) as four_day_avg
from adv_sales
```

```

where extract (month from sales_day) = 5
and   extract (year from sales_day) = 2015
order  by sales day;

```

SALES_DAY	SALES	FOUR_DAY_AVG
01-MAY-15	425	217.50
02-MAY-15	10	253.33
03-MAY-15	325	315.00
04-MAY-15	500	346.25
05-MAY-15	550	343.75
06-MAY-15	0	262.50
07-MAY-15	0	187.50
08-MAY-15	200	162.50
09-MAY-15	450	307.50
10-MAY-15	580	413.75
11-MAY-15	425	482.50
12-MAY-15	475	463.75
13-MAY-15	375	443.75
14-MAY-15	500	500.00
15-MAY-15	650	518.75
16-MAY-15	550	537.50
17-MAY-15	450	537.50
18-MAY-15	500	518.75
19-MAY-15	575	493.75
20-MAY-15	450	506.25
21-MAY-15	500	525.00
22-MAY-15	575	593.75
23-MAY-15	850	606.25
24-MAY-15	500	625.00
25-MAY-15	575	606.25
26-MAY-15	500	537.50
27-MAY-15	575	537.50
28-MAY-15	500	537.50
29-MAY-15	575	556.25
30-MAY-15	575	556.25
31-MAY-15	575	575.00

31 rows selected.

### 3.3. Logical Windowing

The previous windowing clauses have used a certain number of rows. We can also have Oracle calculate which rows fit into our window groups. This is often done with date values where we might want the average of the previous week's sales. In that case the range would be

| Range between Interval '7' day preceding and current row

This syntax requires a datetime field (year, month, day, hour, minute, second)

**Demo 18:** We want a total of the quantity ordered for this month and the previous month. This uses the oe tables

```

with CTE as (
  select order_date, sum(quantity_ordered) As AMount
  from oe_orderheaders
  join oe_orderdetails using (order_id)
  where extract(year from order_date)= 2016
  group by order_date
)
select order_date, amount
, sum(amount) over (
  order by order_date range between

```

```

        interval '1' month preceding and current row
    ) as MonthSum
  from CTE
;

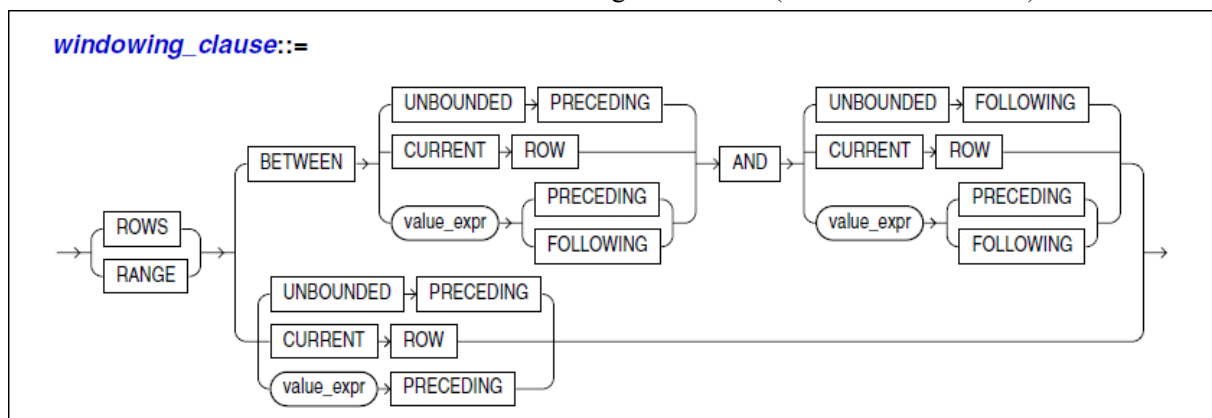
```

ORDER_DATE	AMOUNT	MONTHSUM
02-JAN-16	10	10
03-JAN-16	21	31
04-JAN-16	1	32
05-JAN-16	29	61
07-JAN-16	1	62
11-JAN-16	1	63
12-JAN-16	4	67
15-JAN-16	10	77
23-JAN-16	5	82
26-JAN-16	3	85
27-JAN-16	7	92
31-JAN-16	24	116
01-FEB-16	16	132
02-FEB-16	2	134
03-FEB-16	6	130
12-FEB-16	7	84
01-MAR-16	53	84
05-MAR-16	54	114
07-MAR-16	20	134
08-MAR-16	14	148
09-MAR-16	5	153
04-APR-16	12	105
05-APR-16	4	109
06-APR-16	1	56
07-APR-16	3	59
08-APR-16	4	43
01-MAY-16	86	110
09-MAY-16	4	90
12-MAY-16	15	105

29 rows selected

Again the first few rows of data may not be as valuable, since we do not have a full month of prior data. This is more practical and meaningful if we have a lot of data.

The windowing clause can be quite flexible; this is the model for that clause from the Oracle documentation. Windows can be based on a number of rows on a logical interval (often based on time).



The analytic functions can occur only in the Select list or the Order By clause. Other parts of the query ( join, Where, Group by, Having) are carried out before the analytic functions.

***analytic\_function::=***



***analytic\_clause::=***

